

Assignment 5 — MLOps

ViT Fine-tuning with LoRA & Adversarial Attacks

Tejas Gupta
B22CS093
IIT Jodhpur

Links

- **GitHub:** <https://github.com/TejasGupta-27/MLOps-TejasGupta-B22CS093/tree/Assignment5>
- **WandB (Q1):** <https://wandb.ai/b22cs093-prom-iit-rajasthan/vit-lora-cifar100>
- **WandB (Q2):** <https://wandb.ai/b22cs093-prom-iit-rajasthan/adversarial-cifar10>
- **HuggingFace:** <https://huggingface.co/Tron2703/vit-lora-cifar100>

All experiments were run inside a Docker container on an NVIDIA RTX A5000 (24 GB) GPU using PyTorch with mixed-precision (FP16) training.

1 Q1: ViT-S Fine-tuning with LoRA on CIFAR-100

1.1 Setup

- **Model:** `vit_small_patch16_224` (timm), pre-trained on ImageNet-1K.
- **Dataset:** CIFAR-100 resized to 224×224 , 90/10 train-val split.
- **LoRA:** Injected into Q, K, V attention weights via the PEFT library.
- **Training:** 10 epochs, AdamW ($\text{lr} = 10^{-3}$), cosine annealing, batch size 128.
- **Dropout:** 0.1 for all LoRA experiments.

1.2 Experiment 1 (No LoRA): Head-only Fine-tuning

Trainable Parameters: **38,500** (classification head only).

Table 1: Experiment 1 — No LoRA

Epoch	Train Loss	Val Loss	Train Acc (%)	Val Acc (%)
1	1.0235	0.7360	73.82	79.54
2	0.5860	0.6770	82.76	80.72
3	0.5149	0.6661	84.46	80.68
4	0.4668	0.6640	85.66	81.08
5	0.4353	0.6552	86.70	81.12
6	0.4025	0.6525	87.65	81.24
7	0.3832	0.6436	88.13	81.68
8	0.3649	0.6415	88.84	81.58
9	0.3508	0.6375	89.36	81.38
10	0.3439	0.6379	89.51	81.40

1.3 Experiment 2: Rank=2, Alpha=2

Trainable Parameters: **36,864**.

Table 2: Experiment 2 — LoRA R=2, $\alpha=2$, Dropout=0.1

Epoch	Train Loss	Val Loss	Train Acc (%)	Val Acc (%)
1	3.3436	2.3649	24.61	44.26
2	2.0346	1.8988	54.03	58.12
3	1.7212	1.7152	62.38	62.14
4	1.5637	1.6098	66.47	64.46
5	1.4661	1.5434	68.56	66.14
6	1.3970	1.4898	70.41	67.20
7	1.3501	1.4600	71.54	67.20
8	1.3164	1.4361	72.50	68.60
9	1.2928	1.4330	73.14	68.44
10	1.2820	1.4262	73.24	68.60

1.4 Experiment 3: Rank=2, Alpha=4

Trainable Parameters: **36,864**.

Table 3: Experiment 3 — LoRA R=2, $\alpha=4$, Dropout=0.1

Epoch	Train Loss	Val Loss	Train Acc (%)	Val Acc (%)
1	3.0484	2.2490	32.78	51.02
2	1.8777	1.8658	59.44	60.22
3	1.6193	1.6975	65.23	64.00
4	1.4826	1.5806	68.56	66.54
5	1.3953	1.5206	70.70	68.02
6	1.3238	1.4590	72.45	69.82
7	1.2781	1.4350	73.72	70.24
8	1.2400	1.4100	74.41	70.54
9	1.2160	1.3966	75.04	70.60
10	1.2032	1.3923	75.56	70.86

1.5 Experiment 4: Rank=2, Alpha=8

Trainable Parameters: **36,864**.

Table 4: Experiment 4 — LoRA R=2, $\alpha=8$, Dropout=0.1

Epoch	Train Loss	Val Loss	Train Acc (%)	Val Acc (%)
1	2.9794	2.0348	32.27	54.46
2	1.8107	1.6815	59.51	62.68
3	1.5556	1.5441	65.93	66.14
4	1.4142	1.4437	69.56	68.12
5	1.3280	1.3610	72.00	70.36
6	1.2637	1.3236	73.63	71.36
7	1.2106	1.2973	75.05	71.68
8	1.1703	1.2659	76.14	72.76
9	1.1426	1.2556	76.91	73.10
10	1.1268	1.2400	77.18	73.06

1.6 Experiment 5: Rank=4, Alpha=2

Trainable Parameters: **73,728**.

Table 5: Experiment 5 — LoRA R=4, $\alpha=2$, Dropout=0.1

Epoch	Train Loss	Val Loss	Train Acc (%)	Val Acc (%)
1	3.3300	2.3193	24.78	44.68
2	1.8309	1.7354	57.22	59.68
3	1.4614	1.5082	67.11	66.12
4	1.2622	1.3512	71.93	68.94
5	1.1424	1.2695	74.80	72.14
6	1.0592	1.2207	76.72	72.80
7	1.0024	1.1756	78.21	73.64
8	0.9635	1.1571	79.18	74.38
9	0.9378	1.1415	79.80	74.56
10	0.9253	1.1408	80.21	74.70

1.7 Experiment 6: Rank=4, Alpha=4

Trainable Parameters: **73,728**.

Table 6: Experiment 6 — LoRA R=4, $\alpha=4$, Dropout=0.1

Epoch	Train Loss	Val Loss	Train Acc (%)	Val Acc (%)
1	2.9392	1.8776	33.82	56.52
2	1.5459	1.3755	65.70	70.18
3	1.2223	1.1668	73.58	74.98
4	1.0521	1.0632	77.67	77.52
5	0.9503	0.9848	79.77	79.06
6	0.8804	0.9370	81.56	80.24
7	0.8288	0.9092	82.92	80.62
8	0.7908	0.8898	83.87	80.58
9	0.7689	0.8802	84.26	80.88
10	0.7575	0.8721	84.60	81.24

1.8 Experiment 7: Rank=4, Alpha=8

Trainable Parameters: **73,728**.

Table 7: Experiment 7 — LoRA R=4, $\alpha=8$, Dropout=0.1

Epoch	Train Loss	Val Loss	Train Acc (%)	Val Acc (%)
1	2.6024	1.5670	41.99	66.30
2	1.2907	1.1959	72.12	73.74
3	1.0309	1.0728	77.72	75.56
4	0.9004	0.9566	80.67	78.76
5	0.8125	0.9151	82.69	79.26
6	0.7535	0.8754	83.99	79.46
7	0.7052	0.8534	85.37	80.58
8	0.6671	0.8316	86.48	80.98
9	0.6417	0.8249	87.20	80.66
10	0.6284	0.8192	87.41	80.88

1.9 Experiment 8: Rank=8, Alpha=2

Trainable Parameters: **147,456**.

Table 8: Experiment 8 — LoRA R=8, $\alpha=2$, Dropout=0.1

Epoch	Train Loss	Val Loss	Train Acc (%)	Val Acc (%)
1	3.3195	2.1397	26.03	51.04
2	1.6880	1.4772	61.24	66.08
3	1.2464	1.1762	71.62	73.10
4	1.0178	1.0272	76.71	76.72
5	0.8815	0.9349	80.04	78.60
6	0.7917	0.8800	82.19	79.16
7	0.7352	0.8366	83.47	80.58
8	0.6979	0.8149	84.40	81.20
9	0.6745	0.8052	84.93	81.18
10	0.6623	0.8010	85.18	81.32

1.10 Experiment 9: Rank=8, Alpha=4

Trainable Parameters: **147,456**.

Table 9: Experiment 9 — LoRA R=8, $\alpha=4$, Dropout=0.1

Epoch	Train Loss	Val Loss	Train Acc (%)	Val Acc (%)
1	2.9010	1.7009	34.47	59.58
2	1.3090	1.1719	69.36	72.18
3	0.9583	0.9683	77.68	77.30
4	0.7908	0.8713	81.49	78.88
5	0.6915	0.7906	83.75	81.02
6	0.6207	0.7684	85.61	81.50
7	0.5718	0.7372	86.80	82.06
8	0.5365	0.7271	87.85	82.06
9	0.5146	0.7156	88.42	82.24
10	0.5045	0.7142	88.79	82.54

1.11 Experiment 10: Rank=8, Alpha=8

Trainable Parameters: **147,456**.

Table 10: Experiment 10 — LoRA R=8, $\alpha=8$, Dropout=0.1

Epoch	Train Loss	Val Loss	Train Acc (%)	Val Acc (%)
1	2.6181	1.3861	40.24	68.44
2	1.0836	0.9385	75.07	78.14
3	0.7863	0.7919	81.90	81.60
4	0.6432	0.7200	85.15	82.64
5	0.5538	0.6691	87.31	83.92
6	0.4899	0.6355	89.05	84.56
7	0.4466	0.6161	90.22	84.88
8	0.4138	0.6053	91.17	85.18
9	0.3892	0.6025	91.85	85.08
10	0.3769	0.5999	92.41	85.22

1.12 Test Results Summary

Table 11: Overall Test Accuracy across all configurations

LoRA	Rank	Alpha	Dropout	Test Acc (%)	Trainable Params
Without	—	—	—	81.11	38,500
With	2	2	0.1	68.54	36,864
With	2	4	0.1	70.10	36,864
With	2	8	0.1	72.55	36,864
With	4	2	0.1	74.45	73,728
With	4	4	0.1	80.10	73,728
With	4	8	0.1	80.96	73,728
With	8	2	0.1	81.26	147,456
With	8	4	0.1	82.96	147,456
With	8	8	0.1	85.40	147,456

1.13 Analysis of Q1 Results

Effect of Rank. Increasing the LoRA rank from 2 to 8 consistently improves both validation and test accuracy. At rank=2, the LoRA models (36,864 trainable params) *underperform* the head-only baseline (38,500 params), indicating that the low-rank bottleneck is too restrictive for the QKV adaptation required. At rank=8, LoRA models (147,456 params) comfortably surpass the baseline. Higher rank allows the low-rank matrices to span a richer subspace of weight updates, enabling more expressive fine-tuning of the attention mechanism.

Effect of Alpha. For every fixed rank, raising alpha from 2 to 8 improves performance. Alpha controls the scaling factor α/r applied to the LoRA update $\Delta W = \frac{\alpha}{r}BA$. A larger alpha amplifies the learned perturbation, allowing LoRA to make more meaningful adjustments to the frozen attention weights. The best configuration ($R=8$, $\alpha=8$) achieves **85.40%** test accuracy—4.29 percentage points above the head-only baseline.

Parameter Efficiency. Even the best LoRA model uses only 147,456 trainable parameters, which is 0.68% of the full 21.7M ViT-S model. This demonstrates the strong parameter efficiency of LoRA for transfer learning on CIFAR-100.

Convergence. Low-rank/low-alpha configurations (e.g. $R=2$, $\alpha=2$) converge slowly and plateau at lower accuracy, while $R=8$, $\alpha=8$ converges faster and reaches higher final accuracy, suggesting that the scaling factor plays a critical role in training dynamics.

1.14 Class-wise Test Accuracy Histograms




Figure 1: Class-wise test accuracy for the baseline and best LoRA configuration. The LoRA model shows more uniform accuracy across classes.




Figure 2: Class-wise accuracy for selected LoRA configurations, showing progressive improvement with increasing rank and alpha.

1.15 Gradient Update Graphs on LoRA Weights

Gradient norms on LoRA parameters (lora_A and lora_B matrices across all attention layers) were logged to WandB at every training step. The gradients show a clear decreasing trend over training, consistent with the cosine annealing learning rate schedule. Higher-rank configurations exhibit larger initial gradient magnitudes but converge to similar small values. The full gradient norm plots are available on the WandB dashboard: <https://wandb.ai/b22cs093-prom-iit-rajasthan/vit-lora-cifar100>

1.16 Optuna Hyperparameter Search

Optuna was used to search over rank $\in \{2, 4, 8\}$, alpha $\in \{2, 4, 8\}$, and learning rate (log-uniform in $[10^{-4}, 10^{-2}]$) with 20 trials and median pruning.

Best configuration found:

- Rank = 8, Alpha = 8, LR = 0.00564
- Best Validation Accuracy: **87.56%**

This confirms the grid-search finding that higher rank and alpha are optimal. The Optuna-tuned learning rate (0.00564) is $\sim 5.6\times$ higher than the default (10^{-3}), indicating that LoRA parameters benefit from a more aggressive learning rate since only a small fraction of the network is being updated.

The best Optuna model weights are pushed to GitHub (`weights/best_optuna_model.pth`) and HuggingFace.

2 Q2: Adversarial Attacks using IBM ART

2.1 Q2(i): FGSM Attack — From Scratch vs IBM ART

2.1.1 ResNet18 Training on CIFAR-10

A ResNet18 model (modified for CIFAR-10: 3×3 first conv, removed maxpool) was trained from scratch for 50 epochs using SGD (lr = 0.1, momentum = 0.9, milestones at epochs 20/35/45).

Best Test Accuracy: 93.34% (well above the 72% threshold).

2.1.2 FGSM Implementation

From Scratch: Computed the gradient of the cross-entropy loss w.r.t. the input image, then perturbed each pixel by $\epsilon \cdot \text{sign}(\nabla_x L)$.

IBM ART: Used `FastGradientMethod` from the Adversarial Robustness Toolbox with the same epsilon values.

2.1.3 Perturbation Strength vs Accuracy

Table 12: FGSM Attack: Accuracy at various perturbation strengths

ϵ	Clean Acc (%)	FGSM Scratch (%)	FGSM ART (%)
0.00	93.35	93.35	93.35
0.01	93.35	65.84	70.16
0.02	93.35	44.32	48.81
0.05	93.35	24.93	29.36
0.10	93.35	18.60	22.72
0.15	93.35	16.11	19.71
0.20	93.35	14.37	17.04
0.30	93.35	11.90	13.34




Figure 3: Perturbation strength (ϵ) vs model accuracy under FGSM attack. Both implementations cause catastrophic accuracy drops, with the scratch version being slightly more aggressive.

2.1.4 Visual Comparison



Figure 4: Top: original images. Middle: FGSM from scratch ($\epsilon = 0.1$). Bottom: FGSM via IBM ART ($\epsilon = 0.1$). Red titles indicate misclassification.

2.1.5 Analysis

FGSM is remarkably effective: even at $\epsilon = 0.01$, accuracy drops from 93.35% to 66–70%. At $\epsilon = 0.1$, accuracy plummets below 20%, despite the perturbations being largely imperceptible to humans.

The ART implementation yields 2–5% higher adversarial accuracy compared to the scratch version. This difference arises because ART applies additional numerical safeguards such as gradient clipping and careful handling of floating-point precision, producing slightly less aggressive perturbations.

The key takeaway is that standard DNNs are highly vulnerable to even single-step gradient-based attacks, motivating the need for adversarial training or certified defenses.

2.2 Q2(ii): Adversarial Detection Model

2.2.1 Setup

ResNet-34 binary classifiers were trained to distinguish clean images from adversarial images. Adversarial examples were generated using IBM ART’s PGD and BIM attacks with $\epsilon = 0.15$, step size 0.02, and 20 iterations on 10,000 CIFAR-10 training images. The detector was trained for 25 epochs with Adam ($\text{lr} = 0.001$).

2.2.2 Adversarial Sample Visualization




Figure 5: Clean (top row) vs adversarial (bottom row) images for PGD and BIM attacks ($\epsilon = 0.15$).

2.2.3 Detection Results

Table 13: Adversarial Detection Accuracy (binary: clean vs adversarial)

Attack Type	Detection Accuracy (%)
PGD (Projected Gradient Descent)	99.15
BIM (Basic Iterative Method)	99.90

Both detectors achieve $\gg 70\%$ detection accuracy.

2.2.4 Analysis

PGD and BIM are both iterative extensions of FGSM. PGD applies FGSM repeatedly with random initialization and projects perturbations back onto the ϵ -ball after each step. BIM applies FGSM iteratively with small step sizes without random restarts.

With $\epsilon = 0.15$ and 20 iterations, both attacks produce perturbations that introduce detectable statistical artifacts. The ResNet-34 detector learns to identify subtle patterns in the spatial and frequency domains of adversarial images that deviate from natural image statistics.

BIM achieves slightly higher detectability (99.90% vs 99.15%) because its perturbations follow a more uniform, predictable pattern without random restarts, making them marginally easier to distinguish from clean data.

The near-perfect detection rates demonstrate that adversarial perturbations, while imperceptible to humans, leave identifiable statistical fingerprints that deep learning models can exploit for detection.

Full sample visualizations (10 clean + adversarial pairs for FGSM, PGD, and BIM) are available on WandB: <https://wandb.ai/b22cs093-prom-iit-rajasthan/adversarial-cifar10>

3 Conclusion

1. **LoRA for ViT:** LoRA enables parameter-efficient fine-tuning of ViT-S on CIFAR-100. The best configuration ($R=8$, $\alpha=8$) achieves 85.40% test accuracy using only 0.68% of the model’s parameters, surpassing the head-only baseline by 4.29%.
2. **Rank and Alpha:** Both rank and alpha have a strong positive effect on accuracy. Higher rank increases the expressivity of the low-rank adaptation, while higher alpha amplifies the learned update signal. Optuna confirmed $R=8$, $\alpha=8$ as optimal with a tuned LR of 0.00564.
3. **FGSM Vulnerability:** A single gradient step (FGSM) can reduce a 93% accurate model to below 20% at $\epsilon = 0.1$. Both scratch and ART implementations produce similar results, with ART being marginally less aggressive due to numerical safeguards.
4. **Adversarial Detection:** Deep learning-based detectors (ResNet-34) achieve near-perfect accuracy (99.15% PGD, 99.90% BIM) in distinguishing clean from adversarial images, demonstrating that adversarial perturbations leave detectable statistical fingerprints despite being imperceptible to humans.